

Operating System Overview (part 2)

Operating Systems

João Vicente Ferreira Lima

Universidade Federal de Santa Maria
jvlima@inf.ufsm.br
<http://www.inf.ufsm.br/~jvlima>

2020/2

Outline

1 Introduction to Operating System

2 Unix Overview

- Unix History
- The BSDs

3 Linux Overview

- The GNU project
- Linux kernel

4 Fundamental Concepts

- The Kernel
- The Shell
- Users and Groups
- Files and Directories
- System Calls
- Processes

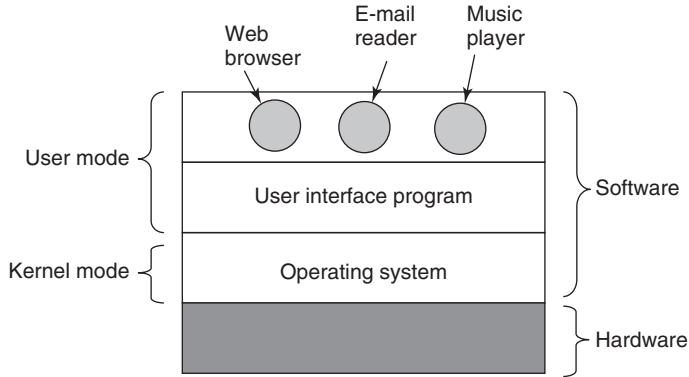
5 FreeBSD

- FreeBSD

Operating System

The term *operating system* is commonly used with two meanings:

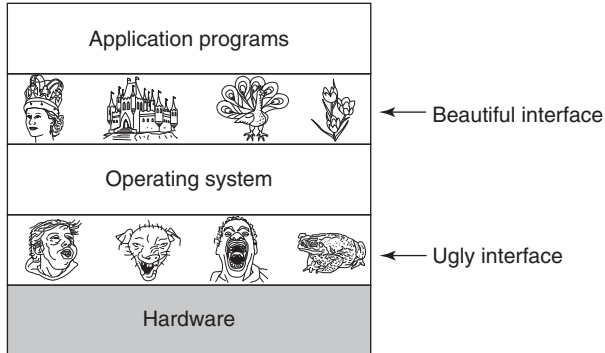
- 1 To denote the entire package: the central software such as command-line interpreters, GUI, file utilities, and editors.
- 2 To the central software that manages and allocates computer resources.



Operating System

The term *operating system* is commonly used with two meanings:

- 1 To denote the entire package: the central software such as command-line interpreters, GUI, file utilities, and editors.
- 2 To the central software that manages and allocates computer resources.



Outline

1 Introduction to Operating System

2 Unix Overview

- Unix History
- The BSDs

3 Linux Overview

- The GNU project
- Linux kernel

4 Fundamental Concepts

- The Kernel
- The Shell
- Users and Groups
- Files and Directories
- System Calls
- Processes


5 FreeBSD

- FreeBSD

Unix history

UNIX was written in assembly by Ken Thompson at Bell Labs (AT&T) in 1969 for a Digital PDP-7 minicomputer. It took several ideas from MULTICS such as:

- tree-structured file system.
- separate program for interpreting commands (shell).
- notion of files as unstructured streams of bytes.



Denis Ritchie designed and implemented C language between 1969-1973 at Bell Labs. By 1973, UNIX was almost totally written in C.

C for system programming

In the 70s, widely used languages were designed to other purposes:

- FORTRAN for mathematical tasks
- COBOL for commercial systems

PDP-7 minicomputer



Unix popularity


At this time, UNIX was popular in universities. The terms of AT&T' agreement with the US government prevented it from selling software. The UNIX university distributions included documentation and the kernel source code (10k lines).

BSD and System V

- **BSD** or *Berkeley Software Distribution* was a UNIX version created when Thompson spent the 1975/1976 academic year as a visiting professor at Berkeley. It included *C shell*, *vi*, *sendmail*, a Pascal compiler, etc.
- **System V** was a comercial version of UNIX when AT&T was permitted to market UNIX.

Other implementations: SunOS or Solaris, Ultrix and OSF/1, IBM's AIX, HP-UX, NeXT's NeXTEP, A/UX for the Apple Macintosh, XENIX.

Unix architecture



- The kernel is relatively small, and controls the hardware resources.
- The interface to the kernel is the system calls. Libraries are built on top of it, but applications are free to use both.
- The shell is a special application that provides an interface for running other applications.

Outline

1 Introduction to Operating System

2 Unix Overview

- Unix History
- The BSDs

3 Linux Overview

- The GNU project
- Linux kernel

4 Fundamental Concepts

- The Kernel
- The Shell
- Users and Groups
- Files and Directories
- System Calls
- Processes

5 FreeBSD

- FreeBSD

The BSDs

It is worth noting that another free UNIX was already available during the early 1990s. Bill and Lynne Jolitz had developed a port of the already mature BSD system (Net/2 from AT&T) for the x86-32, known as 386/BSD. Alternatives from the 386/BSD were:

- NetBSD, FreeBSD, OpenBSD, DragonFly BSD (FreeBSD 4.x)

AT&T and Berkeley lawsuit

In April 1992, USL (AT&T) filed suit against BSDi (Berkeley) and the University of California claiming that the product was still encumbered by proprietary USL source code and trace secrets. A countersuit by the university ensued since USL had not given credit for the BSD code in System V.

The verdict: Berkeley was required to remove 3 of the 18k files from Net/2 release, make some minor changes, and add USL copyright, resulting in the 4.4BSD-Lite. The terms of legal settlement required BSDi, FreeBSD and NetBSD to replace their Net/2 base.

Outline

1 Introduction to Operating System

2 Unix Overview

- Unix History
- The BSDs

3 Linux Overview

- The GNU project
- Linux kernel

4 Fundamental Concepts

- The Kernel
- The Shell
- Users and Groups
- Files and Directories
- System Calls
- Processes

5 FreeBSD

- FreeBSD

The GNU project

Richard Stallman started the GNU project to develop an entire, freely available, UNIX-like system, consisting of software+kernel. Initially they produced a wide range of other programs, but did not produce a working UNIX kernel.

GNU software

GNU C compiler, *bash* shell, GNU C library.

GNU/HURD

GNU/HURD project is a kernel based on the Mach microkernel, but never released.

Outline

1 Introduction to Operating System

2 Unix Overview

- Unix History
- The BSDs

3 Linux Overview

- The GNU project
- Linux kernel

4 Fundamental Concepts

- The Kernel
- The Shell
- Users and Groups
- Files and Directories
- System Calls
- Processes

5 FreeBSD

- FreeBSD

The Linux kernel

In 1991, Linus Torvalds was inspired to write an OS for his Intel 80386 PC. In the course of his studies, he had come into contact with Minix, a small UNIX-like operating system developed by Andrew Tanenbaum.

The first version was released by March 1994.

Linux distributions

- Slackware in 1993
- Debian soon after
- SUSE and Red Hat followed
- Ubuntu in 2004

IBM commercial: <https://www.youtube.com/watch?v=x7ozaFbqg00>

MEMORABLE LINUX MILESTONES

CELEBRATING 20 YEARS OF LINUX

LINUX TORVALDS
POSTS FAMOUS
MESSAGE - "HELLO
EVERYBODY OUT
THERE..." - AND
RELEASES FIRST
LINUX CODE



1991

SLACKWARE
BECOMES FIRST
WIDELY ADOPTED
DISTRIBUTION



1993

TECH GIANTS
BEGIN ANNOUNCING
PLATFORM SUPPORT
FOR LINUX



1998

IBM RUNS
FAMOUS LINUX
AD DURING THE
SUPERBOWL



2003

THE LINUX
FOUNDATION IS
FORMED TO PROMOTE
PROTECT AND
STANDARDIZE LINUX
LINUX IS A FELLOW



2007

LINUX TURNS 20
AND POWERS THE
WORLD'S
SUPERCOMPUTERS,
STOCK EXCHANGES,
PHONES, ATMS,
HEALTHCARE
RECORDS,
SMART GRIDS, THE
LIST GOES ON



2011



LINUX LICENSES
LINUX UNDER
THE GPL, AN
IMPORTANT
DECISION THAT
WILL CONTRIBUTE
TO ITS SUCCESS IN
THE COMING YEARS



LINUX VISITS
AQUARIUM, GETS
BIT BY A PENGUIN
AND CHOOSES
IT AS LINUX MASCOT



RED HAT
GOES PUBLIC



LINUX APPEARS ON
THE COVER OF
BUSINESSWEEK, WITH
A STORY THAT HAILS
LINUX AS A
BUSINESS SUCCESS



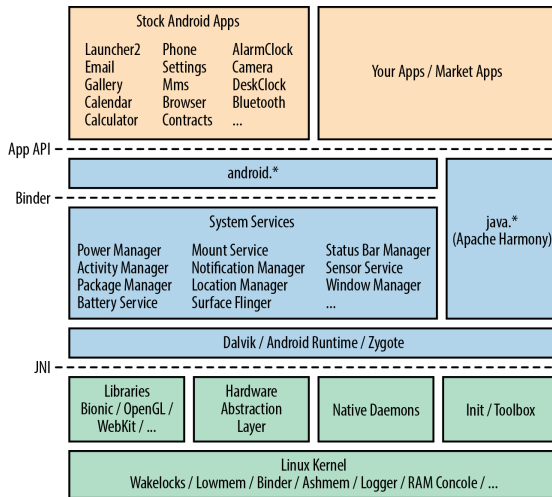
THE LINUX-BASED
ANDROID OS
OUTSHIPS ALL OTHER
SMARTPHONE OSes
IN THE U.S. AND
CLIMBS TO
DOMINANCE



THE
LINUX
FOUNDATION
<http://www.linuxfoundation.org/>

Copyright © 2011 Linux Foundation. All rights reserved. Linux is a registered trademark of Linus Torvalds. Slackware is a registered trademark of Slackware Linux, Inc. Red Hat is a registered trademark of Red Hat, Inc. Android is a trademark of Google Inc. Use of this trademark is subject to Google Permissions. BusinessWeek is a trademark of BLOOMBERG L.P.

Android



Outline

1 Introduction to Operating System

2 Unix Overview

- Unix History
- The BSDs

3 Linux Overview

- The GNU project
- Linux kernel

4 Fundamental Concepts

- The Kernel
- The Shell
- Users and Groups
- Files and Directories
- System Calls
- Processes

5 FreeBSD

- FreeBSD

The Kernel role

Tasks performed by the kernel:

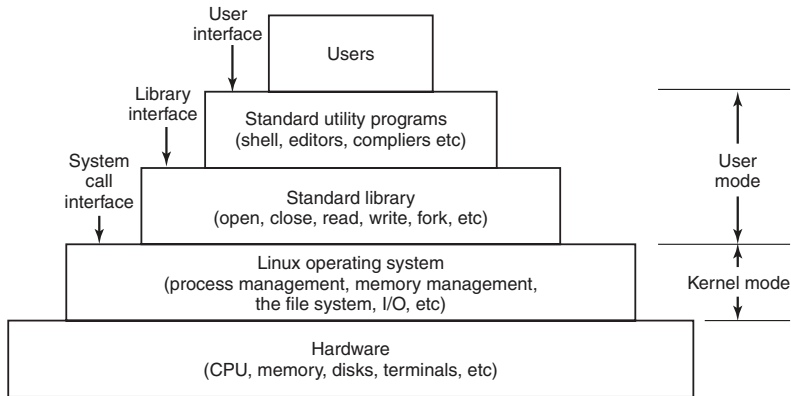
- **Process scheduling:** Linux is a *preemptive multitasking* SO.
- **Memory management:** virtual memory management.
- **File system:** file I/O.
- **Creation and termination of processes:** allocating resources.
- **Access to devices:** interface between programs and devices.
- **Networking:** transmit and receive network messages on behalf of user processes.
- **System call API:** entry points to request various tasks.

Kernel mode vs user mode

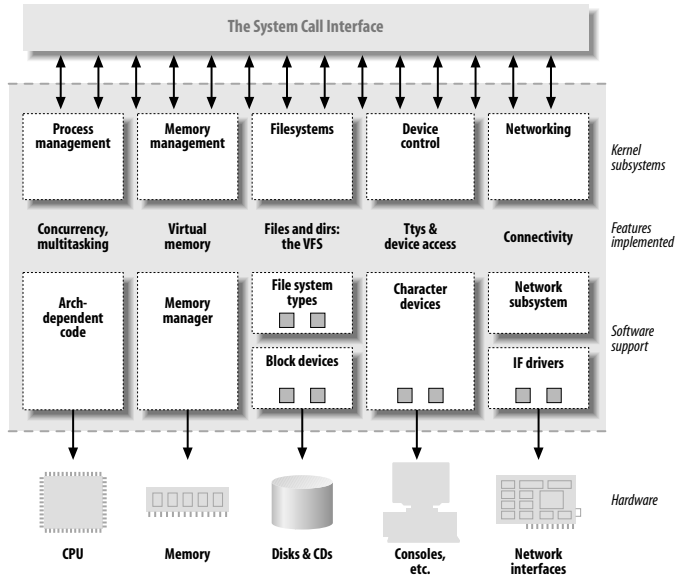
Typically there are at least two different modes: *user mode* and *kernel mode* (or *supervisor mode*). When running in user mode, the CPU can access only memory that is marked as begin in user space.

Certain operations can be performed only in kernel mode.

Linux system layers



Linux kernel layers



Outline

1 Introduction to Operating System

2 Unix Overview

- Unix History
- The BSDs

3 Linux Overview

- The GNU project
- Linux kernel

4 Fundamental Concepts

- The Kernel
- **The Shell**
- Users and Groups
- Files and Directories
- System Calls
- Processes

5 FreeBSD

- FreeBSD

The Shell

The *shell* or *command interpreter* is a special-purpose program designed to read commands and execute them. On some OS the shell is an integral part of the kernel, on UNIX is a user process. Examples:

- **Bourne shell (sh)**: the oldest of the widely used shells, written by Steve Bourne.
- **C shell (csh)**: similar to C on control-flow constructs.
- **Korn shell (ksh)**: successor of the Bourne shell.
- **Bourne again shell (bash)**: GNU Bourne shell with features from C and Korn shells.

The shells are designed not only for interactive use, but also for the interpretation of *shell scripts* (text with commands). For this purpose, each has features such as: variables, loop and conditional statements, I/O commands, and functions.

Outline

1 Introduction to Operating System

2 Unix Overview

- Unix History
- The BSDs

3 Linux Overview

- The GNU project
- Linux kernel

4 Fundamental Concepts

- The Kernel
- The Shell
- **Users and Groups**
- Files and Directories
- System Calls
- Processes

5 FreeBSD

- FreeBSD

Users and groups

Users: every user has a unique *login name* (username) and *user ID* (UID). There are defined in the system *password file* `/etc/passwd` with:

- **Group ID:** the GID of first group the user is a member.
- **Home directory:** initial user directory.
- **Login shell:** name of the program to interpret user commands.

The password is recorded in encrypted form in the *shadow password file*.

Groups: users are organized in *groups* The *group file* `/etc/group` has:

- **Group name:** unique group name.
- **Group ID (GID):** numeric ID.
- **User list:** list of login names of users who are members.

Superuser: or *root* has special privileges within the system.

Outline

1 Introduction to Operating System

2 Unix Overview

- Unix History
- The BSDs

3 Linux Overview

- The GNU project
- Linux kernel


4 Fundamental Concepts

- The Kernel
- The Shell
- Users and Groups
- **Files and Directories**
- System Calls
- Processes

5 FreeBSD

- FreeBSD

Files and directories



File types: each file is marked with a *type*. Ordinary files are *regular* or *plain* files. Other types are devices, pipes, sockets, directories, and symbolic links.

Directories: a *directory* is a special file whose contents take the form of a table of filenames. Each directory entry is a *link* or *hard link*. Every directory has at least two entries:

- **.** (**dot**) which is a link to the directory itself
- **..** (**dot-dot**) which is a link to its *parent directory*.

Symbolic links: while a hard link points to the data itself, a *soft link* only provides the text of an *actual* file name.

Filename: On most Linux systems, filenames can be up to 255 characters long. We should avoid non-portable characters in filenames (`[-._a-zA-Z0-9]`).

Pathname: is a string with an optional initial slash (`/`) followed by a series of filenames separated by slashes.

- **Absolute pathname:** begin with a slash (`/`). Ex: `/usr/include`.
- **Relative pathname:** relative to a process's current working directory. Ex: `../bin/exec`.

File ownership and permissions

Each file has an associated UID and GID that define its owner and group's owner. The system divides users into three categories of file access:

- **owner** (*user*) of the file
- **group** users who are members of the group
- **other** the rest of the world

Three permission bits may be set for each user:

- **read**
- **write**
- **execute**

In directories, these permissions have different meanings:

- **read** allows to list the contents (filenames)
- **write** add/delete/change contents
- **execute** (or *search*) access to files within the directory

File ownership and permissions

Binary	Symbolic	Allowed file accesses
111000000	rwX-----	Owner can read, write, and execute
111111000	rwXrwX---	Owner and group can read, write, and execute
110100000	rw-r-----	Owner can read and write; group can read
110100100	rw-r--r--	Owner can read and write; all others can read
111101101	rwXr-Xr-X	Owner can do everything, rest can read and execute
000000000	-----	Nobody has any access
000000111	-----rwx	Only outsiders have access (strange, but legal)

One of the features of UNIX systems is the concept of *universality of I/O*, i.e, system calls perform I/O on all types of files, including devices.

File descriptors - the I/O system calls refer to open files by a *file descriptor*, a non-negative integer. Normally, a process inherits three open file descriptors when it is started by the shell:

- **0** (*standard input*) from it takes the input.
- **1** (*standard output*) to which the process writes its output.
- **2** (*standard error*) to which the process writes error messages.

File descriptions

```
char c = fgetc( 0 );  
fprintf( 1 , "%c\n", c);
```

The *stdio* library

```
char c = fgetc(stdin);  
fprintf(stdout, "%c\n", c);
```


Outline

1 Introduction to Operating System

2 Unix Overview

- Unix History
- The BSDs

3 Linux Overview

- The GNU project
- Linux kernel

4 Fundamental Concepts

- The Kernel
- The Shell
- Users and Groups
- Files and Directories
- **System Calls**
- Processes

5 FreeBSD

- FreeBSD

System Calls

The interface between OS and programs. Most calls are performed by system libraries (`libc`).


File management

Call	Description
<code>fd = open(file, how, ...)</code>	Open a file for reading, writing, or both
<code>s = close(fd)</code>	Close an open file
<code>n = read(fd, buffer, nbytes)</code>	Read data from a file into a buffer
<code>n = write(fd, buffer, nbytes)</code>	Write data from a buffer into a file
<code>position = lseek(fd, offset, whence)</code>	Move the file pointer
<code>s = stat(name, &buf)</code>	Get a file's status information

Process management

Call	Description
<code>pid = fork()</code>	Create a child process identical to the parent
<code>pid = waitpid(pid, &statloc, options)</code>	Wait for a child to terminate
<code>s = execve(name, argv, environp)</code>	Replace a process' core image
<code>exit(status)</code>	Terminate process execution and return status

Call to read()



Outline

1 Introduction to Operating System

2 Unix Overview

- Unix History
- The BSDs

3 Linux Overview

- The GNU project
- Linux kernel

4 Fundamental Concepts

- The Kernel
- The Shell
- Users and Groups
- Files and Directories
- System Calls
- Processes

5 FreeBSD

- FreeBSD

Process definition

A **process** is an instance of an executing program. The kernel:


- ① loads the code into virtual memory
- ② allocates space for program variables
- ③ sets up kernel bookkeeping data structures to record various information)

Process memory layout

A process is divided into the following parts, known as *segments*:

- **text** - the instructions of the program.
- **data** - the static variables.
- **heap** - an area from which programs dynamically allocate memory.
- **stack** - piece of memory that grows and shrinks as functions are called and return, used to allocate storage for local variables and function call linkage information.

Process memory layout



Process creation

The `fork()` system call creates a new process. The process that calls `fork()` is the *parent process*, and the new *child process*. Additional steps:

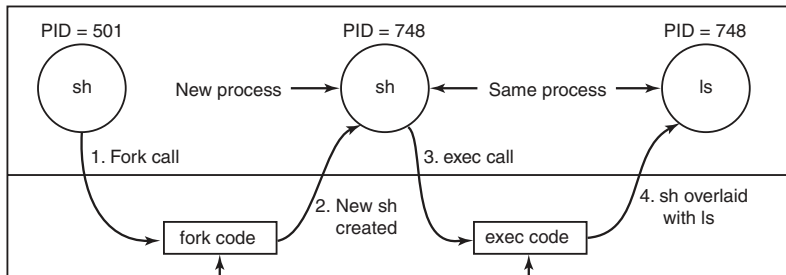
- the kernel creates the child by making a duplicate of the parent
- the child inherits copies of the parent's data, stack, and heap segments.

The child process goes on either to execute a different set of functions, or to use the `execve()` system call to load and execute a entirely new program.

Process identifier

Each process has a unique *process identifier* (PID), and a *parent process identifier* (PPID).

A ls from the shell



Allocate child's task structure
Fill child's task structure from parent
Allocate child's stack and user area
Fill child's user area from parent
Allocate PID for child
Set up child to share parent's text
Copy page tables for data and stack
Set up sharing of open files
Copy parent's registers to child

Find the executable program
Verify the execute permission
Read and verify the header
Copy arguments, environ to kernel
Free the old address space
Allocate new address space
Copy arguments, environ to stack
Reset signals
Initialize registers

Process tree

```
init-+-acpid
      |-atd
      |-cron
      |-cups-browsed
      |-cupsd
      |-exim4
      |-6*[getty]
      |-irqbalance
      |-portmap
      |-redis-server---2*[{redis-server}]
      |-rsyslogd-+-{in:imklog}
      |           |-{in:imuxsock}
      |           ‘-{rs:main Q:Reg}
      |-screen---bash
      |-sshd-+-sshd---sshd---bash---pstree
      |       ‘-2*[sshd---sshd---bash]
```

Special-purpose processes

The *init* process - when booting the system, the kernel creates a special process called *init*, the parent of all processes, which is derived from the program `/sbin/init`. All programs from the system are created either by *init* or by one of its descendants.

The *init* process can't be killed, and terminates only when the system is shut down.

Daemon processes - a *daemon* is a special-purpose process that is created and handled by the system, but:

- It is long-lived, in general started at system boot and remains in existence until the system is shut down.
- It runs in background, and has no controlling terminal from which it can read input or to which it can write output.

Outline

1 Introduction to Operating System

2 Unix Overview

- Unix History
- The BSDs

3 Linux Overview

- The GNU project
- Linux kernel

4 Fundamental Concepts

- The Kernel
- The Shell
- Users and Groups
- Files and Directories
- System Calls
- Processes

5 FreeBSD

- FreeBSD

FreeBSD overview

FreeBSD is a 4.4BSD-Lite based operating system for Intel and AMD64 computers. Software can be installed using:

- **pkg** is the binary package management for FreeBSD.
- **Ports Collection** which is a set of Makefiles, patches, and description files stored in `/usr/ports`.

Systems using FreeBSD

- **Playstation 4** - Sony's PlayStation 4 Is Running Modified FreeBSD 9
- **Apple OS X** and some of **iOS**
- Apache, Whatsapp, Netflix, etc



FreeBSD®

